

QUORDA

Neutral procurement clearing for autonomous buyers and sellers.

Agent Tank Product & Builder Compendium | Internal score 94/100

QUORDA - Product Pack

Tagline: Neutral procurement clearing for autonomous buyers and sellers.

Category: Agent-to-agent procurement

Internal winner-grade score: 94/100

Product thesis

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

Why this pack exists

This folder is a complete conception-to-build baseline for the Agent Tank build window. It is intended to let one builder/team move from pitch to a functioning frontend + Intelligent Contract product without rediscovering product decisions.

Canonical product boundary

QUORDA decides who should win before a deal is formed. PRAEST governs whether an already-selected service provider later fulfilled its obligations.

Non-negotiables

- Not a generic procurement SaaS.
- Do not let an LLM evaluate hard numeric constraints.
- No escrow dependency.
- Hackathon bids are synthetic/public; privacy architecture is documented for production.

Documentation map

01 executive brief; 02 pitch; 03 problem/market; 04 vision/scope; 05 PRD; 06 TRD; 07 architecture; 08 Intelligent Contract; 09 evidence/decision design; 10 UX; 11 privacy/security; 12 testing; 13 two-week build; 14 SDLC; 15 demo/submission; 16 GTM/metrics; 17 risks/ADRs/roadmap; 18 research/differentiation; AGENTS.md builder operating rules.

Executive Brief

The product

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

The real problem

Procurement agents can already draft RFQs and compare quotes, but cross-organisation autonomous procurement still depends on one buyer-controlled model to interpret ambiguous bids, non-price trade-offs and public supplier evidence. That model is not neutral to the seller agents whose bids it ranks.

Primary users

- buyer agents
- supplier agents
- open B2B agent marketplaces
- SMEs automating sourcing
- protocols that need impartial award selection

Core loop

1. Buyer agent posts an RFQ with hard constraints, soft priorities and evidence policy.
2. Seller agents submit offers with structured commercial fields and source/evidence URLs.
3. Deterministic logic eliminates bids that violate hard constraints.
4. GenLayer ranks the remaining bids against the published natural-language priorities and independently checkable evidence.
5. The winning bid receives an AWARDED commitment; ties or insufficient evidence return NEEDS_CLARIFICATION instead of hallucinating certainty.
6. Payment or contract-formation adapters can act on the award without QUORDA becoming an escrow product.

Why GenLayer belongs here

Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example "best support coverage without sacrificing delivery certainty" - against evidence that all validators can inspect.

Consequential state change

one bid becomes AWARDED, the RFQ closes, and downstream purchase/contract formation can reference the award receipt.

Market wedge

Procurement automation exists, but impartial machine-to-machine clearing is still open territory. QUORDA can sit between A2A discovery and payment, turning competitive agent offers into a neutral, auditable award.

Business model

Clearing fee per RFQ/award; marketplace API; enterprise procurement policy packs. No native token and no forced escrow.

Hackathon success definition

A stranger should be able to use the live app, trigger one clean case and one adversarial/ambiguous case, inspect the Intelligent Contract transaction, understand why validator judgment was necessary, and verify that the resulting verdict changed product state.

Agent Tank Pitch Pack

One-line pitch

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

X pitch draft

AI agents can already talk, transact and use tools. But agent-to-agent procurement still needs trustworthy judgment at the moment it becomes consequential.

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

Hard constraints stay deterministic. GenLayer handles only the ambiguous decision that independent validators need to judge from shared evidence. The verdict changes real product state, not a dashboard score.

Built for the agentic economy as a standalone product, with GenLayer exactly where neutral judgment is necessary.

90-second video script

AI agents are becoming economic actors. Procurement agents can already draft RFQs and compare quotes, but cross-organisation autonomous procurement still depends on one buyer-controlled model to interpret ambiguous bids, non-price trade-offs and public supplier evidence.

That is why we are building QUORDA.

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

We do not send everything to AI. Hard rules stay deterministic. Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example "best support coverage without sacrificing delivery certainty" - against evidence that all validators can inspect.

When consensus is reached, something real changes: one bid becomes AWARDED, the RFQ closes, and downstream purchase/contract formation can reference the award receipt.

A buyer agent needs an API provider under a hard monthly budget and latency ceiling, while preferring strong refund terms, support coverage and documented uptime. Three seller agents bid. One is cheapest but violates latency, one passes hard checks but has weak support, and one wins after GenLayer evaluates the declared trade-offs and public evidence.

QUORDA is not a blockchain wrapper around an AI demo. It is infrastructure for a machine economy where autonomous actors need neutral, evidence-backed decisions they can all verify.

Panel memory hook

QUORDA: Neutral procurement clearing for autonomous buyers and sellers.

What makes the pitch shortlistable

- A concrete agentic-economy failure, not "AI + blockchain."
- A narrow decision GenLayer is uniquely suited to judge.
- A real state consequence after consensus.
- A demo that visibly fails under naive automation and succeeds with neutral adjudication.
- A credible path to becoming a standalone product after the hackathon.

Problem and Market Thesis

Problem statement

Procurement agents can already draft RFQs and compare quotes, but cross-organisation autonomous procurement still depends on one buyer-controlled model to interpret ambiguous bids, non-price trade-offs and public supplier evidence. That model is not neutral to the seller agents whose bids it ranks.

Why this becomes worse with autonomous agents

Humans can pause, ask a lawyer, call a supplier, inspect a diff or return to checkout. Autonomous systems are valuable precisely because they act without waiting for a person. That removes the informal human judgment layer and creates demand for explicit, machine-readable decisions that counterparties can verify independently.

Existing stack and missing layer

A2A handles communication/task exchange; x402 handles machine-native payment; ERC-8004 addresses identity/reputation/validation hooks; AP2 addresses verifiable payment mandates. None makes every ambiguous business decision disappear. QUORDA occupies one sharply defined gap rather than rebuilding those standards.

Beachhead users

- buyer agents
- supplier agents
- open B2B agent marketplaces
- SMEs automating sourcing
- protocols that need impartial award selection

Why now

- Agent interoperability and payments are becoming standards rather than demos.
- More agents are being given economic and operational authority.
- The cost of an incorrect autonomous decision scales with that authority.
- GenLayer makes neutral consensus on unstructured evidence an application primitive.

Market thesis

Procurement automation exists, but impartial machine-to-machine clearing is still open territory. QUORDA can sit between A2A discovery and payment, turning competitive agent offers into a neutral, auditable award.

Product Vision and Scope

Vision

Make agent-to-agent procurement safe enough for autonomous actors to use without placing a single company, model or operator in the position of final judge.

Product principles

7. Deterministic before non-deterministic.
8. Judgment must be narrow, explicit and reviewable.
9. Every verdict must change consequential product state.
10. Evidence must be independently inspectable by validators.
11. Uncertainty is a valid output; never hallucinate certainty.
12. No native token is required for product-market fit.
13. Store commitments and receipts onchain; minimise sensitive plaintext.

In scope for two weeks

- Polished responsive frontend.
- Injected wallet flow.
- One canonical GenLayer Intelligent Contract.
- Live/synthetic evidence inputs appropriate to the product.
- Clean case, rejection case and uncertainty/edge case.
- Explorer-verifiable transactions and receipts.
- Public deployment and reproducible demo data.

Explicitly out of scope

- Custodial funds or regulated escrow.
- Native token issuance.
- Broad multi-chain production infrastructure.
- Private enterprise data connectors in the hackathon MVP.
- Claims of legal, security or financial guarantees beyond the contract's defined decision.

Product boundary vs PRAEST

QUORDA decides who should win before a deal is formed. PRAEST governs whether an already-selected service provider later fulfilled its obligations.

Product Requirements Document (PRD)

Objective

Ship a credible MVP of QUORDA that demonstrates a complete decision lifecycle from user/agent intent through neutral adjudication to consequential state.

Personas

- ****Buyer Agents**** - needs a neutral, portable decision without trusting one counterparty model.
- ****Supplier Agents**** - needs a neutral, portable decision without trusting one counterparty model.
- ****Open B2B Agent Marketplaces**** - needs a neutral, portable decision without trusting one counterparty model.
- ****Smes Automating Sourcing**** - needs a neutral, portable decision without trusting one counterparty model.
- ****Protocols That Need Impartial Award Selection**** - needs a neutral, portable decision without trusting one counterparty model.

Primary user story

As an autonomous or human-controlled actor, I want to submit a decision request with explicit rules and evidence so that a neutral validator network can determine the ambiguous part and produce an outcome downstream systems can trust.

Functional requirements

- FR-01: create/register the product-specific object.
- FR-02: hash and display immutable policy/criteria before judgment.
- FR-03: submit evidence or public evidence references.
- FR-04: run deterministic validation before invoking GenLayer reasoning.
- FR-05: request consensus judgment with a bounded output schema.
- FR-06: expose pending/accepted/final/appealable status truthfully.
- FR-07: persist a portable receipt with request hash, evidence hash, verdict, reasoning summary, timestamps and transaction identifiers.
- FR-08: prevent prohibited state transitions.
- FR-09: support at least one negative case and one uncertainty case.
- FR-10: provide an explorer link for every consequential transaction.

Product-specific lifecycle

14. Buyer agent posts an RFQ with hard constraints, soft priorities and evidence policy.
15. Seller agents submit offers with structured commercial fields and source/evidence URLs.
16. Deterministic logic eliminates bids that violate hard constraints.
17. GenLayer ranks the remaining bids against the published natural-language priorities and independently checkable evidence.
18. The winning bid receives an AWARDED commitment; ties or insufficient evidence return NEEDS_CLARIFICATION instead of hallucinating certainty.
19. Payment or contract-formation adapters can act on the award without QUORDA becoming an escrow product.

Non-functional requirements

- No mock claim presented as live data.
- Frontend + contract first; no required application backend.
- Clear loading/finality/error states.
- Mobile usable at 390px width.
- All contract calls fail closed on malformed evidence or impossible state.
- Reproducible fixtures for judges.

Success metrics

- 100% of demo verdicts trace to an onchain transaction.
- 0 UI actions offered when the contract would reject the transition.
- At least 3 end-to-end scenarios pass: clean, reject, ambiguous/uncertain.
- At least 1 scenario proves why deterministic-only logic is insufficient.

Technical Requirements Document (TRD)

Reference stack

- Frontend: Next.js + TypeScript.
- Wallet: injected EIP-1193 provider only.
- Chain interaction: viem or genlayer-js where appropriate.
- Intelligent Contract: Python / GenLayer.
- Hosting: Vercel-compatible static/serverless frontend features only; core correctness cannot depend on a proprietary backend.
- Tests: frontend unit/component tests plus contract schema/state-transition tests.

Contract invariants

- A request cannot be judged before required deterministic checks pass.
- Verdict schema is finite and machine-readable.
- Evidence identifiers are immutable once judgment begins.
- Accepted is not presented as final when an appeal/finality window still applies.
- Downstream consequential actions consume the recorded verdict, not a frontend re-computation.

Product-specific methods

- ``create_rfq(rfq_hash, deadline, hard_criteria, soft_policy_hash)``
- ``submit_bid(rfq_id, bid_hash, bid_capsule_uri)``
- ``close_bidding(rfq_id)``
- ``filter_hard_constraints(rfq_id)``
- ``judge_award(rfq_id)``
- ``accept_award(rfq_id, bid_id)``
- ``get_award_receipt(rfq_id)``

States

- ``OPEN``
- ``BIDDING_CLOSED``
- ``FILTERED``
- ``UNDER_JUDGMENT``
- ``AWARDED``
- ``NO_VALID_BID``
- ``NEEDS_CLARIFICATION``
- ``CANCELLED``

Evidence types

- RFQ specification
- hard and soft criteria
- seller bid payloads
- public supplier evidence URLs
- optional ERC-8004 identity/capability references

Non-deterministic boundary

Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example “best support coverage without sacrificing delivery certainty” - against evidence that all validators can inspect.

Performance target


The UI should acknowledge transaction submission immediately, display consensus/finality progress truthfully, and avoid pretending GenLayer judgment is an instant synchronous API call.

Failure handling

- Source unavailable -> SOURCE_UNAVAILABLE / retry path.
- Evidence hash mismatch -> reject before judgment.
- Malformed model output -> fail closed; do not mutate consequential state.
- Ambiguous evidence -> explicit uncertainty/clarification state where supported.
- User rejects wallet transaction -> local cancelled state only; no fabricated chain status.

System Architecture

Human / Agent / Organisation



Component responsibilities

Frontend: compose requests, hash local artefacts, call wallet/contract, display authoritative chain state and receipts.

Deterministic contract logic: identity/authorization, deadlines, numeric constraints, state transitions, hash integrity and simple filters.

GenLayer non-deterministic logic: only the product-specific semantic judgment.

Evidence sources: public URLs or minimal capsules validators can independently inspect.

Adapters: optional interoperability, never required for the core demo.

Core product judgment

Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example “best support coverage without sacrificing delivery certainty” - against evidence that all validators can inspect.

Consequential state

one bid becomes AWARDED, the RFQ closes, and downstream purchase/contract formation can reference the award receipt.

Privacy boundary

Bids can contain sensitive commercial terms. The hackathon demo uses public/synthetic bids. Production design stores bid commitments onchain and encrypts full bids off-chain, with selective adjudication capsules. Avoid exposing real supplier pricing on a public testnet.

Intelligent Contract Specification

Contract responsibility

The Intelligent Contract is the canonical state machine and adjudicator for QUORDA. The frontend may format data but cannot manufacture a verdict or bypass state rules.

Public interface

- ``create_rfq(rfq_hash, deadline, hard_criteria, soft_policy_hash)``
- ``submit_bid(rfq_id, bid_hash, bid_capsule_uri)``
- ``close_bidding(rfq_id)``
- ``filter_hard_constraints(rfq_id)``
- ``judge_award(rfq_id)``
- ``accept_award(rfq_id, bid_id)``
- ``get_award_receipt(rfq_id)``

State machine

OPEN -> BIDDING_CLOSED -> FILTERED -> UNDER_JUDGMENT -> AWARDED -> NO_VALID_BID -> NEEDS_CLARIFICATION -> CANCELLED

The literal implementation may branch rather than follow a single line; this diagram shows the lifecycle vocabulary. Every method must define allowed source states and destination states explicitly.

Judgment schema

```
{
  "decision": "PRODUCT_SPECIFIC_ENUM",
  "confidence_band": "high|medium|low",
  "material_findings": ["short factual finding"],
  "unresolved_questions": ["question requiring clarification"],
  "evidence_refs": ["source identifier"],
  "policy_version": "hash-or-version"
}
```

Consensus prompt design

- State the exact decision question in one sentence.
- Separate immutable hard facts from natural-language policy.
- Require citation/reference to supplied evidence identifiers.
- Forbid introduction of external assumptions not supported by evidence.
- Require uncertainty when evidence is missing or contradictory.
- Keep output small and strongly structured for equivalence checking.

Security invariants

- Verdict can only mutate state from a judgment-ready state.
- Replayed evidence cannot overwrite a completed object.
- Caller permissions are explicit.
- Hashes bind policy/evidence versions to the receipt.
- No settlement/release occurs from frontend text alone.

Evidence and Decision Design

Evidence set

- RFQ specification
- hard and soft criteria
- seller bid payloads
- public supplier evidence URLs
- optional ERC-8004 identity/capability references

Evidence quality rules

20. Prefer primary/current sources over commentary.
21. Bind URLs to retrieved timestamps and content hashes where practical.
22. Make missing evidence visible rather than silently substituting assumptions.
23. Keep deterministic facts outside the LLM prompt where code can enforce them.
24. Store a concise evidence manifest so a judge can reproduce the decision.

Decision boundary

Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example “best support coverage without sacrificing delivery certainty” - against evidence that all validators can inspect.

Output discipline

The verdict must use a product-specific finite enum, list the material reasons that affected the result, and expose unresolved issues. Long free-form essays are not contract state.

Adversarial cases to test

- Source disappears between request and validation.
- One source contradicts another.
- Evidence is irrelevant but persuasive-looking.
- User supplies a prompt-injection string inside evidence.
- A party changes a referenced document after commitment.
- Validator output is semantically correct but phrased differently.

Receipt

Each receipt should contain request/object id, policy hash, evidence manifest hash, decision, timestamp/block/tx identifiers, and finality status.

UX and User Flows

UX goal

Make a hard protocol concept understandable in under one minute: what is being decided, which facts are deterministic, what evidence validators inspect, what consensus returned, and what changed because of it.

Required screens

25. Landing / product thesis.
26. Workspace dashboard showing objects and state.
27. Create/request flow.
28. Evidence/criteria review before wallet submission.
29. Judgment progress with truthful chain status.
30. Result detail with reasons, evidence, transaction and receipt.
31. Demo scenarios page with reset/replay guidance.

Golden path

32. Buyer agent posts an RFQ with hard constraints, soft priorities and evidence policy.
33. Seller agents submit offers with structured commercial fields and source/evidence URLs.
34. Deterministic logic eliminates bids that violate hard constraints.
35. GenLayer ranks the remaining bids against the published natural-language priorities and independently checkable evidence.
36. The winning bid receives an AWARDED commitment; ties or insufficient evidence return NEEDS_CLARIFICATION instead of hallucinating certainty.
37. Payment or contract-formation adapters can act on the award without QUORDA becoming an escrow product.

Result language

Never show “AI says...” as the authority. Show “validator consensus” or the product verdict, with a clear link to evidence and transaction state.

Accessibility

Keyboard navigable controls; visible focus states; status not conveyed by colour alone; readable contrast; responsive tables/cards; error messages tied to fields.

Privacy, Security and Threat Model

Privacy posture

Bids can contain sensitive commercial terms. The hackathon demo uses public/synthetic bids. Production design stores bid commitments onchain and encrypts full bids off-chain, with selective adjudication capsules. Avoid exposing real supplier pricing on a public testnet.

Assets to protect

- User intent and private commercial details.
- Evidence integrity.
- Contract state integrity.
- Wallet authorization boundaries.
- Honest representation of finality.

Threats and controls

| Threat | Control |

| --- | --- |

| Evidence tampering | Hash evidence/policy before judgment; reject mismatch. |

| Prompt injection in web evidence | Treat evidence as data, constrain prompt and output schema, require independent validation. |

| Frontend bypass | Contract enforces state transitions and consequences. |

| Sensitive data leakage | Commitments/minimum disclosure; synthetic/public data in MVP. |

| Stale web source | Timestamp/hash source snapshot; expose source age. |

| Ambiguous result | Explicit uncertainty state; fail closed. |

| Finality confusion | Separate submitted/accepted/final/appealable states in UI. |

| Wallet overreach | Injected wallet only; no secret/private key collection. |

Production privacy direction

Encrypted off-chain payloads + onchain commitments + minimal adjudication capsules. Optional PET/ZK adapters may be explored later, but the core product must not depend on a commercial privacy chain agreement to function.

Testing and Acceptance Specification

Test layers

- Pure deterministic functions: hashes, numeric filters, state guards.
- Contract transition tests: every valid and invalid state change.
- Judgment parser/schema tests: malformed, missing and contradictory outputs fail closed.
- Frontend tests: button availability exactly matches contract rules.
- End-to-end: wallet -> transaction -> consensus -> receipt -> UI.

Mandatory scenarios

Clean pass: A buyer agent needs an API provider under a hard monthly budget and latency ceiling, while preferring strong refund terms, support coverage and documented uptime.

Negative case: same workflow with a material violation that must not pass.

Uncertainty case: remove or contradict a key evidence source and require an uncertain/clarification outcome instead of a fabricated answer.

Acceptance gates

- All contract state-transition tests pass.
- No action button exists in a contract-invalid state.
- Every verdict contains policy/evidence hashes.
- Explorer transaction link resolves.
- At least one real non-deterministic GenLayer call is exercised in the demo.
- No mock dataset is presented as live external evidence.
- README reproduction steps work on a clean clone.

Two-Week Agent Tank Build Plan

Build window: 3 September 2026, 12:00 UTC to 17 September 2026.

Date	Milestone	Exit gate
---	---	---
Sep 3	Freeze product contract, schemas, state machine	No unresolved core scope
Sep 4	Implement deterministic contract skeleton	State tests green
Sep 5	Implement GenLayer judgment path	Structured verdict returned
Sep 6	Adversarial prompt/evidence hardening	Fail-closed tests green
Sep 7	Frontend shell + injected wallet	Connect/read/write works
Sep 8	Primary create/request flow	Object visible from chain state
Sep 9	Result/receipt/finality UI	Tx + evidence trace visible
Sep 10	Negative and uncertainty fixtures	All 3 scenarios reproducible
Sep 11	Public deployment integration	Live app talks to deployed contract
Sep 12	UX polish + mobile pass	No broken critical flow
Sep 13	Full automated test sweep	Release branch green
Sep 14	Demo rehearsal + bug burn-down	3 consecutive clean demos
Sep 15	Docs, architecture, screenshots, explorer evidence	Submission pack complete
Sep 16	External dry run by teammate	Fresh user completes demo unaided
Sep 17	Final smoke test and submit before deadline	Immutable submission evidence saved

Scope protection

If behind schedule, cut adapters and analytics first. Never cut the consequential contract state, evidence trace, negative case or reproducibility.

SDLC and Engineering Standard

Branch discipline

main is releasable. Feature branches are short-lived. Merge only after tests and a human-readable change note.

Definition of done

- Requirement implemented and mapped to an acceptance test.
- Invalid state transitions tested.
- Typecheck/lint/tests green.
- No secrets or private keys committed.
- UI reflects contract lifecycle exactly.
- Documentation updated when public interface or state changes.

Change control

Architecture decisions that alter custody, privacy, adjudication authority, evidence sources, finality or economic consequence require an ADR in 17_RISKS_ADRS_ROADMAP.md.

Release evidence

Record commit SHA, deployed contract address, deployment tx, live URL, test counts and the three canonical demo transaction IDs.

Engineering rule

The product must remain understandable without hidden infrastructure. A judge should be able to inspect the frontend repository, Intelligent Contract and onchain evidence and see where every important decision occurs.

Demo and Submission Playbook

2-minute demo

38. State the problem in 15 seconds.
39. Show the immutable criteria/policy before judgment.
40. Run the adversarial case first so the danger is obvious.
41. Open the GenLayer transaction/receipt and show why it was rejected or held.
42. Run the corrected case and show the consequential state transition.
43. Close with where this plugs into the agentic stack and why it can be a standalone product.

Canonical demo

A buyer agent needs an API provider under a hard monthly budget and latency ceiling, while preferring strong refund terms, support coverage and documented uptime. Three seller agents bid. One is cheapest but violates latency, one passes hard checks but has weak support, and one wins after GenLayer evaluates the declared trade-offs and public evidence.

Proof checklist

- GitHub repository.
- Live deployment.
- Contract address and explorer link.
- Clean-case transaction.
- Negative-case transaction.
- Uncertainty/edge transaction if supported.
- Test command and result.
- Architecture diagram.
- README with 5-minute reproduction path.

Pitch line to avoid

Do not say “GenLayer makes it trustless” without showing the exact decision and state consequence. The demo must make the trust boundary visible.

GTM, Business Model and Metrics

Beachhead

Procurement automation exists, but impartial machine-to-machine clearing is still open territory. QUORDA can sit between A2A discovery and payment, turning competitive agent offers into a neutral, auditable award.

Pricing thesis

Clearing fee per RFQ/award; marketplace API; enterprise procurement policy packs. No native token and no forced escrow.

No-token rule

The product does not need a native token. GEN is a network/protocol asset; commercial value can remain asset-agnostic. A token is not a substitute for demand.

North-star metric

Consequential decisions successfully completed and consumed by a downstream workflow.

Operating metrics

- Decision requests per week.
- Pass/reject/uncertain distribution.
- Median time to accepted/final outcome.
- Appeal rate and overturn rate.
- Evidence-source failure rate.
- Percentage of verdicts actually consumed downstream.
- Repeat integrator usage.

GTM path

44. Hackathon live demo with public SDK-style interface.
45. Integrate one adjacent standard/ecosystem rather than building a closed marketplace.
46. Recruit 3-5 design partners from the narrow beachhead.
47. Measure repeated consequential use before expanding scope.
48. Add enterprise privacy/indexing only after the core verdict loop is proven.

Risks, ADRs and Roadmap

Principal risks

| Risk | Mitigation |

| --- | --- |

| R1 - Too generic | Keep one narrow adjudication question and one beachhead. |

| R2 - Forced GenLayer use | Deterministic logic stays deterministic; GenLayer only handles semantic judgment. |

| R3 - Privacy leakage | Public/synthetic MVP; commitments + encrypted capsules for production. |

| R4 - Hallucinated certainty | Explicit uncertain/clarification output and fail-closed state. |

| R5 - Weak evidence | Primary sources, timestamps, hashes and reproducible manifests. |

| R6 - Demo-only product | Business model and downstream integration must survive after hackathon. |

ADR-001: No native token

Decision: no product token in MVP.

Reason: product utility must stand on decisions, not emissions.

ADR-002: Frontend + contract first

Decision: no required central application backend for canonical correctness.

Reason: simplifies reproducibility and keeps authority visible.

ADR-003: Minimum disclosure

Decision: public chain receives commitments/receipts; sensitive payloads remain offchain where possible.

Reason: GenLayer currently does not make public-chain data private by default.

Roadmap

Phase 0 - Agent Tank: one decision type, one strong evidence pattern, one polished interface.

Phase 1: SDK + adapter for the most relevant adjacent standard.

Phase 2: encrypted evidence capsules, indexing, richer policy/version management.

Phase 3: ecosystem integrations and production economics.

Research and Differentiation

Core differentiation

QUORDA decides who should win before a deal is formed. PRAEST governs whether an already-selected service provider later fulfilled its obligations.

Why this is not a generic GenLayer demo

Price ceilings and deadlines are deterministic. The GenLayer-native part is interpreting non-price priorities and heterogeneous terms - for example "best support coverage without sacrificing delivery certainty" - against evidence that all validators can inspect.

Competitive/standards boundary

- A2A: communication and task orchestration, not this product's neutral consequential judgment.
- ERC-8004: identity/reputation/validation primitives; QUORDA may integrate where relevant but does not replace the standard.
- x402: payment transport, not the judgment layer.
- AP2: payment authorization/mandates, not a replacement for this product's domain unless explicitly integrated.
- Internet Court / PRAEST: post-agreement disputes/service assurance, whereas QUORDA is scoped to agent-to-agent procurement.

Research sources

| Source | URL | Why it matters |

| --- | --- | --- |

| Agent Tank Hackathon | <https://portal.genlayer.foundation/agent-tank/> | Original agentic-economy ideas; pre-season pitch; build window 3-17 September 2026; ten ideas shortlisted. |

| GenLayer - What is GenLayer | <https://docs.genlayer.com/understand-genlayer-protocol/what-is-genlayer> | GenLayer is an adjudication layer for decisions requiring language, evidence and judgment. |

| GenLayer - When to Use GenLayer | <https://docs.genlayer.com/developers/intelligent-contracts/when-to-use-genlayer> | Strong fit requires a real consequence, judgment, independently checkable evidence and neutral consensus. |

| GenLayer - Optimistic Democracy | <https://docs.genlayer.com/understand-genlayer-protocol/optimistic-democracy-how-genlayer-works> | Independent AI validators reach consensus; appeals can move accepted outcomes toward finality. |

| A2A Protocol | <https://a2a-protocol.org/dev/specification/> | A2A handles agent discovery, communication and task lifecycle, but does not itself provide neutral adjudication. |

| ERC-8004 | <https://eips.ethereum.org/EIPS/eip-8004> | Identity, reputation and validation registries for trustless agents; advertised capabilities are not automatically guaranteed. |

| x402 | <https://docs.x402.org/introduction> | Internet-native payment standard for agents, APIs and services. |

| Google AP2 | <https://cloud.google.com/blog/products/ai-machine-learning/announcing-agents-to-payments-ap2-protocol> | Signed mandates create verifiable user intent and payment authorization for agentic commerce. |

| AgentClear | <https://agentclear.co.uk/> | Existing governed procurement reinforces demand but is deterministic and organisation-centric; QUORDA differentiates through open cross-agent neutral consensus on ambiguous trade-offs. |

| A2A Protocol | <https://a2a-protocol.org/dev/specification/> | A2A supplies communication/task interoperability; QUORDA supplies neutral award adjudication. |

Name note

The product name has received a preliminary web collision sweep for obvious direct conflicts. This is not legal trademark clearance and should be revisited before commercial launch.

AGENTS.md - QUORDA

You are building QUORDA, not a generic AI/blockchain demo.

Product contract

QUORDA is a neutral clearing layer where a buyer agent publishes an RFQ, seller agents submit complex offers, and GenLayer reaches consensus on which surviving bid best satisfies the buyer's declared trade-offs before an award is committed.

Authority boundary

- Deterministic rules: code.
- Ambiguous product-specific judgment: GenLayer Intelligent Contract.
- Consequence: one bid becomes AWARDED, the RFQ closes, and downstream purchase/contract formation can reference the award receipt.
- Frontend never invents chain truth.

Do not

- Not a generic procurement SaaS.
- Do not let an LLM evaluate hard numeric constraints.
- No escrow dependency.
- Hackathon bids are synthetic/public; privacy architecture is documented for production.
- Do not add a native token.
- Do not add a custodial backend.
- Do not use MetaMask Snaps; injected wallet only.
- Do not present mocked data as live.
- Do not expose private keys or sensitive evidence.

Before every merge

49. Run typecheck/lint/tests.
50. Confirm UI buttons match contract-valid transitions.
51. Re-run clean, reject and uncertainty fixtures.
52. Update docs when method/state/evidence schemas change.
53. Record deployment/test evidence for submission.